

Enterprise AI Platform Operations Policy

Document ID: IT-OPS-001 | Effective: June 27, 2026 | Version: 1.0 | Classification: Internal

APPROVED — ACTIVE

1. Purpose & Scope

This policy defines the responsibilities, standards, and procedures for the deployment, maintenance, monitoring, and availability of the TechCorp Enterprise AI Platform. It establishes operational expectations for the IT Department as the platform operator and sets the service-level objectives that every consuming department can rely upon.

This policy applies to all IT Department personnel responsible for platform operations, including infrastructure engineers, site reliability engineers, and operations support staff. It also applies to any third-party managed service providers contracted to operate platform components, and to **Engineering** and **AI Department** personnel who deploy or modify platform services.

2. AI Platform Integration

The IT Department operates the Enterprise AI Platform using a layered operations model monitored continuously by the platform's own AI-driven observability suite, with oversight from the **AI Department** for model health and the **Security Department** for operational security posture. Core operational integrations include:

- AI-Driven Infrastructure Monitoring** — The platform self-monitors compute, storage, and network health using RAG-powered anomaly detection; alerts are triaged by IT operations engineers and escalated per the incident severity matrix co-maintained with **Security**.
- Automated Deployment Pipelines** — CI/CD pipelines managed by IT in coordination with **Engineering**, with AI-assisted canary analysis and automated rollback triggers gated by pre-deployment safety checks defined by the **AI Department**.
- Capacity Planning & Scaling Engine** — AI forecasting of platform resource demand across departments, generating provisioning recommendations reviewed monthly by IT capacity planners and the **Finance Department** for budget alignment.

- **Service Health Dashboard** — Real-time availability and latency metrics for every platform service, accessible to all department leads, with automated status-page updates triggered by the AI monitoring suite.

Connected Departments: AI (model health & safety deployment gates) • Engineering (CI/CD pipeline & infrastructure-as-code) • Security (operational security monitoring & incident escalation) • Finance (infrastructure cost governance) • Customer Support (platform uptime for AI support tools)

3. Permitted Operations Activities

Activity	AI Role	Human Oversight Required
Platform health monitoring & alerting	Anomaly detection & alert generation	IT engineer triages and acknowledges every P1 alert within 15 minutes
Automated service deployment	Canary analysis & rollback trigger	Release manager approves production deployments; AI rollbacks are reviewed post-hoc
Capacity forecasting & provisioning	Demand prediction & recommendation	IT capacity planner approves provisioning changes above threshold
Patch management & vulnerability remediation	Vulnerability-to-patch mapping	IT engineer validates patch compatibility in staging before production rollout
Backup & disaster recovery testing	DR test scheduling & result analysis	IT DR coordinator reviews results and signs off on RPO/RTO compliance
Log aggregation & retention management	Automated collection & indexing	IT admin confirms retention policy compliance quarterly with Legal

4. Prohibited Activities

The following are strictly prohibited under this policy:

- **Production deployment outside the approved CI/CD pipeline.** All platform changes must flow through the automated deployment pipeline with required test gates. Manual production changes require a break-glass procedure logged and reviewed within 24 hours by IT leadership and Security.
- **Disabling platform monitoring without authorization.** AI-driven monitoring must remain active for all production services. Silencing alerts or disabling monitors requires IT Manager approval and is limited to defined maintenance windows.

- **Exceeding approved maintenance windows without stakeholder notification.**
Planned maintenance exceeding the published window must be communicated to all affected department leads through the Service Health Dashboard at least 2 hours before the window expires.
- **Running platform components past end-of-life or without active security patches.**
All platform infrastructure must be patched within the SLA defined in the vulnerability management policy co-owned with Security.
- **Storing platform operational data outside approved logging infrastructure.** All logs, metrics, and operational data must flow into the centralized observability platform; local-only logs on production systems are prohibited.

5. Availability & Service-Level Objectives

All platform operations must comply with TechCorp's **AI Safety Guidelines** (AI Department) and the **Data Protection Standards** (Security Department). Availability commitments include:

- The Enterprise AI Platform core services target 99.9% uptime measured monthly, excluding planned maintenance windows published 7 days in advance through the Service Health Dashboard.
- Critical path services (RAG pipeline, model inference endpoints, authentication service) target 99.95% uptime with a recovery time objective of 15 minutes for severity-1 incidents.
- Disaster recovery testing is conducted quarterly with a recovery point objective of 1 hour and recovery time objective of 4 hours for full platform restoration.
- Department-specific AI modules inherit the availability tier of their consuming department as defined in the department's AI policy; IT maintains a tiered support matrix aligned with business criticality ratings from Operations.
- All availability data is reported monthly to department heads through the platform's automated SLA reporting module, with violation summaries reviewed by IT leadership and the affected department lead.

6. Operational Quality Standards

IT operations personnel must adhere to TechCorp's **Prompt Engineering Standards** when configuring AI-driven operations tooling. Operations-specific requirements include:

- Monitoring alert thresholds must be tuned per service using the AI-driven threshold recommendation tool, with quarterly review by the service owner to minimize alert fatigue.
- AI-generated incident summaries must be reviewed and annotated by the responding engineer before the incident is closed; unannotated AI summaries are

not accepted as the incident record.

- Capacity forecast prompts must include historical usage data, planned feature releases from Engineering, and seasonal demand patterns from consuming departments to produce calibrated projections.
- All operations runbooks must be version-controlled and tested during each disaster recovery exercise; AI-generated runbook drafts must be reviewed and approved by the IT operations lead before use.

7. Roles & Responsibilities

Role	Responsibility
VP of IT	Overall platform operations ownership; availability SLA accountability; budget approval for infrastructure scaling; escalation point for severity-1 incidents.
IT Operations Manager	Day-to-day operations oversight; maintenance window approval; monitoring coverage management; runbook governance; incident response coordination.
Site Reliability Engineers	Platform health monitoring and alert response; deployment pipeline operation; capacity management; performance tuning; incident root-cause analysis.
Infrastructure Engineers	Compute, storage, and network provisioning; patch management; backup and DR execution; platform component lifecycle management.
AI Department	Model health monitoring within the platform; safety deployment gates for model updates; AI monitoring tool accuracy review; prompt standard updates for operations AI.
Engineering Department	Application-level service ownership; CI/CD pipeline maintenance; code deployment coordination with IT; infrastructure-as-code contributions.
Security Department	Operational security monitoring; vulnerability scanning and remediation SLA enforcement; incident escalation and joint response; access audit coordination.
Finance Department	Infrastructure cost review and budget alignment; capacity provisioning financial approval above defined thresholds.

8. Incident Response

Any of the following constitutes an operational incident and must be reported within the timeframe specified:

- Core platform service outage affecting any department — IT must detect via AI monitoring within 2 minutes and begin response within 5 minutes; affected department leads notified within 15 minutes.

- Data loss or corruption in platform storage systems — report within 30 minutes to IT Operations Manager, Security, and the affected department lead; initiate DR procedures immediately.
- Unauthorized production change or deployment — report within 1 hour to IT Operations Manager and Security; reverse the change and conduct a post-incident review within 48 hours.
- Monitoring system failure resulting in a visibility gap — report within 30 minutes to IT Operations Manager; restore monitoring coverage with manual checks active until automated monitoring is restored.
- Capacity exhaustion approaching service degradation — report within 1 hour to IT capacity planning and Finance; execute emergency provisioning per the approved capacity runbook.

All operational incidents are logged in the platform's incident management module. Severity-1 and severity-2 incidents require a post-incident review within 5 business days involving IT, the affected department, Security, and AI (if model-related). Incident metrics are reported monthly to department heads and quarterly to the board.

9. Training & Compliance

- All IT operations personnel must complete the Enterprise AI Platform Operations certification within 60 days of onboarding and recertify annually.
- Site reliability engineers must complete the AI-Driven Observability workshop covering monitoring configuration, alert tuning, and AI incident summary review standards.
- All personnel with production deployment access must complete the Change Management & Deployment Safety training before receiving pipeline approval privileges.
- Quarterly operational readiness reviews are conducted by the IT Operations Manager with participation from Security and Engineering; gaps are tracked to remediation in the platform's compliance module.
- Non-compliance with this policy is subject to the progressive discipline framework defined in the Employee Handbook and enforced by HR.

10. Review & Revision

This policy is reviewed every 6 months by the VP of IT, IT Operations Manager, and lead site reliability engineer, with input from the AI Department (for model-related operations), Engineering Department (for deployment pipeline coordination), Security Department (for operational security controls), and Finance Department (for infrastructure cost governance). Revisions are versioned, approved by the VP of IT,

and synchronized with updates to the TechCorp AI Safety Guidelines and the Security Incident Response Plan. Emergency revisions triggered by significant platform architecture changes or post-incident findings follow an accelerated 14-day review cycle.

Referenced Policies: AI Safety Guidelines (AI Department) • Prompt Engineering Standards (AI Department) • Data Protection Standards (Security Department) • Security Incident Response Plan (Security Department) • IT Access Management Policy (IT Department) • Enterprise Systems Integration Policy (IT Department) • Employee Handbook (HR Department)