

AI Content Governance & Brand Policy

Document ID: MKT-CONT-001 | Effective: June 27, 2026 | Version: 1.0 | Classification: Internal

APPROVED — ACTIVE

1. Purpose & Scope

This policy defines the requirements for AI-generated marketing content across all TechCorp channels, establishing mandatory human review gates, brand consistency standards, and accountability for every piece of public-facing content. The core principle: AI accelerates content creation, but humans own the final output and the brand.

This policy applies to all Marketing Department personnel who create, edit, or publish content, including content marketers, brand managers, social media managers, and creative teams. It also applies to external agencies and contractors producing content on TechCorp's behalf, and to any other department generating customer-facing content. It operates alongside the **Enterprise Intellectual Property & Data Usage Policy (LEG-CORP-003)** for content ownership and the **Sales AI Governance Policy (SAL-AI-001)** for sales collateral alignment.

2. AI Platform Integration

The Marketing Department uses the Enterprise AI Platform's content studio, with brand compliance enforcement integrated with the **Legal Department's** IP registry and the **AI Department's** safety review tools. Core integrations include:

- AI Content Studio** — Multi-format content generation (copy, social posts, email, ad creative, blog drafts) with built-in brand voice, tone, and style enforcement, producing first drafts for human review with all AI-generated content watermarked as draft until approved.
- Brand Compliance Checker** — AI-driven review of all content against TechCorp's brand guidelines, trademark usage rules, and competitive claim standards, flagging violations before human review with references to the specific guideline breached.
- Human Review Gateway** — Mandatory approval workflow where every piece of AI-generated public-facing content must be reviewed and approved by a qualified marketing professional before publication; the gateway is technically enforced with publication blocked until approval is recorded.

- **Content Attribution Ledger** — Immutable record of AI model version, prompt, human reviewer, and approval timestamp for every published piece of AI-assisted content, enabling full auditability and coordinated with the **Legal Department** for IP and regulatory compliance.

Connected Departments: Legal (brand IP, claims substantiation & regulatory review) • Sales (collateral consistency & competitive messaging) • AI (content model safety & brand prompt governance) • IT (content platform & publication infrastructure) • Customer Support (aligned messaging in support content)

3. Permitted AI-Assisted Content Activities

Activity	AI Role	Human Oversight Required
Copywriting & content drafting	First-draft generation in brand voice	Content marketer reviews, edits, and approves final copy before publication
Brand compliance checking	Automated guideline violation flagging	Brand manager reviews all flags; approves or overrides each finding
Multi-channel content adaptation	Format transformation across channels	Channel manager reviews adapted content for platform-specific fit
SEO & keyword optimization	Keyword analysis & content scoring	SEO specialist validates recommendations; content marketer applies selectively
Image & creative asset generation	Visual concept generation & variation	Creative lead reviews all AI-generated visuals; approves for brand alignment
Content performance prediction	Engagement & conversion forecasting	Content strategist validates predictions against campaign goals

4. Prohibited Activities

The following are strictly prohibited under this policy:

- **Publishing AI-generated content without human review.** Every piece of public-facing content — copy, images, video, social posts, ads, emails — must pass through the Human Review Gateway with documented approval before publication. The gateway is technically enforced; bypassing it is a policy violation.
- **AI generation of unsubstantiated claims.** AI-generated content must not include product capability claims, performance benchmarks, or competitive

comparisons that have not been verified against approved source material from the competitive intelligence repository maintained by Marketing and vetted by **Legal**.

- **AI impersonation of executives or spokespersons.** AI may draft content attributed to named individuals only with that individual's explicit review and approval. AI must never generate content in the voice of a TechCorp executive without their consent and approval of the specific output.
- **Using AI-generated content in regulated communications without Legal review.** Earnings-related content, investor communications, regulatory disclosures, and crisis communications must be drafted by humans and reviewed by Legal — AI may assist with research but not generate first drafts for these categories.
- **Removing or altering AI content attribution metadata.** The Content Attribution Ledger record linking each published piece to its AI model version, prompt, and human reviewer must remain intact and auditable.

5. Brand Standards & Content Integrity

All AI-generated content must comply with TechCorp's **AI Safety Guidelines** (AI Department) and the **Enterprise Intellectual Property & Data Usage Policy (LEG-CORP-003)**. Brand requirements include:

- The AI Content Studio is configured with the current approved brand guidelines, including voice and tone parameters, visual identity specifications, trademark usage rules, and prohibited terminology lists maintained by the Brand team.
- AI-generated content is classified as draft by default in the content management system; the draft classification is removed only upon human approval through the review gateway.
- Every published piece of AI-assisted content carries attribution metadata recording the AI model version, prompt template reference, primary human author, and approving human reviewer for full traceability.
- Brand guideline updates must be reflected in the AI Content Studio configuration within 5 business days of approval; the Brand Manager and AI Department jointly validate that the updated configuration produces compliant outputs.
- All AI-generated content is stored with version history; the AI draft, human-edited version, and final approved version are all preserved for audit and creative process improvement.

6. Content Quality Standards

Marketing personnel must adhere to TechCorp's **Prompt Engineering Standards** when using the AI Content Studio. Content-specific requirements include:

- Content generation prompts must include target audience, channel, content type, brand voice parameters, and any mandatory inclusions (product names, legal disclaimers, CTAs) for the AI to produce usable first drafts.
- Brand compliance prompts must reference the current approved brand guideline version; prompts using outdated guidelines produce unreliable compliance results.
- Multi-channel adaptation prompts must specify the target channel's format constraints, audience expectations, and character limits; the AI must adapt the message without altering the core narrative or claims.
- All content prompt templates must be version-controlled, approved by the Content Strategy Lead, and stored in the department's shared prompt library.
- Quarterly content quality audits sample AI-assisted published content for brand compliance, claim accuracy, and audience engagement; compliance below 98% triggers a prompt template and review process review.

7. Roles & Responsibilities

Role	Responsibility
Chief Marketing Officer (CMO)	Overall content governance ownership; final authority on brand voice and public messaging; approves exceptions to human review requirement (none routine); brand crisis communication leadership.
Brand Manager	Brand guideline maintenance and AI Content Studio configuration; Brand Compliance Checker rule management; trademark usage enforcement; brand consistency audit across all channels.
Content Strategy Lead	Content prompt template governance; editorial calendar alignment with AI production capacity; content quality audit program; Human Review Gateway process oversight.
Content Marketers	AI-assisted content drafting and editing; human review and approval of AI-generated content; prompt usage per approved templates; content performance feedback loop to AI Content Studio.
Creative Lead	AI-generated visual asset review and approval; visual brand consistency enforcement; creative direction for AI visual generation parameters.
Legal Department	Claims substantiation review; trademark and IP compliance in AI-generated content; regulatory communication review; Content Attribution Ledger audit for IP protection.
Sales Department	Collateral consistency feedback; competitive claim validation against field experience; AI-generated sales content alignment with Marketing messaging.
IT Department	Content management and publication platform infrastructure; Human Review Gateway technical enforcement; Content Attribution Ledger storage; access provisioning.
AI Department	Content generation model safety; brand voice model accuracy; content model bias monitoring; content AI prompt standard governance.

8. Incident Response

Any of the following constitutes a content governance incident and must be reported within the timeframe specified:

- AI-generated content published without human review — report within 1 hour to the Content Strategy Lead and Brand Manager; content immediately

reviewed; if non-compliant, retracted or corrected within 2 hours; gateway bypass root cause investigated within 24 hours.

- AI-generated content containing false, misleading, or unsubstantiated claims reaching the public — report within 1 hour to Brand Manager, CMO, and Legal; content retracted immediately; correction published; regulatory implications assessed by Legal.
- AI-generated content causing brand reputation harm — report within 30 minutes to CMO, Brand Manager, and PR lead; crisis communication protocol activated; content retracted; public response coordinated with Legal.
- Content Attribution Ledger tampering or data loss — report within 2 hours to IT, Content Strategy Lead, and Legal; affected records identified; ledger integrity restored; tampering investigated by Security.
- AI Content Studio producing systematically off-brand outputs — Brand Manager and AI Department notified within 4 hours; model configuration reviewed; affected content in pipeline flagged for heightened human review.

All content incidents are logged in the platform's incident management module. Incidents involving public-facing content errors or brand harm require a post-incident review involving Marketing, Legal, and AI Department within 5 business days. Regulatory implications are assessed by Legal within 24 hours.

9. Training & Compliance

- All Marketing personnel must complete the AI Content Governance & Brand Compliance training within 30 days of hire and annually thereafter.
- Content creators must complete the AI-Assisted Content Creation workshop covering prompt engineering for brand voice, human review responsibilities, and prohibited use boundaries before receiving AI Content Studio access.
- Brand and creative personnel must complete the AI Visual Content & Brand Integrity module covering AI-generated image review standards, trademark usage, and visual consistency requirements.
- Quarterly content audits sample published AI-assisted content for compliance; findings are reviewed by the Content Strategy Lead and Brand Manager; audit results reported to the CMO.
- Non-compliance with this policy, particularly publishing AI content without human review or generating unsubstantiated claims, is subject to immediate content tool access revocation and the progressive discipline framework administered by HR.

10. Review & Revision

This policy is reviewed every 6 months by the CMO, Brand Manager, and Content Strategy Lead, with input from the Legal Department (for claims substantiation and regulatory compliance), Sales Department (for content alignment with field needs), IT Department (for platform and gateway infrastructure), and AI Department (for content model safety and brand voice accuracy). Revisions are versioned, approved by the CMO, and synchronized with updates to the Enterprise Intellectual Property & Data Usage Policy, the TechCorp AI Safety Guidelines, and the Prompt Engineering Standards. Emergency revisions triggered by content incidents or brand reputation events follow an accelerated 14-day review cycle.

Referenced Policies: Enterprise Intellectual Property & Data Usage Policy (LEG-CORP-003) • Sales AI Governance Policy (SAL-AI-001) • TechCorp AI Safety Guidelines (AI Department) • Prompt Engineering Standards (AI Department) • Marketing Data Governance & Personalization Policy (MKT-DATA-001) • Brand Reputation & Sentiment Monitoring Policy (MKT-REP-001)